

KARTIK SRINIVAS

kartiksrinivas008@gmail.com ♦ GH:kartiksrinivas007 ♦ LinkedIn ♦ +1(609)-781-8511

EDUCATION

Carnegie Mellon University - MS, Machine Learning (Systems and Theory concentration) GPA: 3.97/4.00 , Coursework: Deep RL, ML Systems, Database Systems, Advanced Stat Theory, PGM	Pittsburgh, PA December 2025
Indian Institute of Technology - B.Tech Computer Science with a Minor in Artificial Intelligence GPA: 9.56/10.00 , Graduated with Honors , an Institute Silver Medal and Triple Dean's List (2021-2024)	Hyderabad, TN July 2024

EXPERIENCE

Research Scientist Waymo LLC Post-training and Reinforcement Learning	Present Mountain View, CA
Research Intern Waymo LLC - Developed post-training methods for Waymo foundation models that improved collision-safety - Built new model attention encoders and launched training jobs for billion-parameter autoregressive models Mined high quality datasets from the driving logs with $< 0.1\%$ selectivity for improving open-loop collision avoidance - Collaborated with predictive planner team to start improvement of on-board Waymo planning models Awarded a peer-bonus for writing efficient TensorProto parsers for a separate infrastructure team within Waymo	Summer 2025 Mountain View, CA
Machine Learning Research Intern UBC Vancouver & Vector Institute for AI - Pioneered decentralized federated mutual learning, enabling privacy-preserving, public-data-free learning Improved state of the art accuracies by 8 % by training 50 models across clients with heterogeneous data - Invented a distributed training algorithm using latent alignment and deep differentially private image generation - Collaborated with Meta AI to write a paper which was accepted to ICML 2024	Summer 2023 Vancouver, BC

PUBLICATIONS

Exact Unlearning via Model Merging at Scale <i>K Kuo, A Setlur, Kartik Srinivas, A Raghunathan, V Smith</i>	ICLR 2025
Overcoming Data and Model Heterogeneities in Decentralized Federated Learning <i>CY Huang, Kartik Srinivas, X Zhang, X Li</i>	ICML 2024
Unlearning Zero-Shot Transfer in Domain Adaptation Through Adversarial Optimization <i>A Devalapally, P Jain, Kartik Srinivas, VN Balasubramanian</i>	CVPR 2026

SKILLS

Frameworks: JAX, Tensorflow, Pytorch, Triton, Numpy, Scikit, Pandas, Wandb, Tensorboard, SLURM
Programming Languages: C++, Python, C, JavaScript, Julia, Matlab, RISC-V Assembly, x86-Assembly

PROJECTS

Triton Kernels for H-Nets and & Mamba - 2 (Mamba-2) - Note: This is present work at Dr. Albert Gu's Lab - Developed GPU Kernels for the Exponential Moving average layer in H-nets beating Mamba-2 GPU Kernel latency - Optimized memory accesses, block sizes and tiling to improve throughput of memory-bound sub-kernels - Implemented alternative parallel prefix scan algorithms for chunked EMA recurrences in triton	Present
DOSD: Decoupled Online Speculative Decoding (Poster) - Developed a novel speculative-decoding algorithm that beats state of the art transformer inference latency - Asynchronously trained a draft-model copy on mistakes made during speculative decoding on another GPU - Co-created a novel weight serialization protocol and utilized MPI for sending weights of draft-model across GPUs	Spring 2025

AWARDS & SERVICE

Awards: MITACS Research Scholar, Shastri-Indo-Canadian Institute Scholar, Google Research week awardee (150 students selected all over India), Institute Academic Excellence (Thrice), Institute Silver Medal, Microsoft Engage Hackathon Awardee
Teaching: Computability Theory, Convex Optimization, Deep Learning, Tensor Analysis & Introduction to Programming